

```

# Standard libraries
import os
import random
from tqdm.notebook import tqdm

# Data manipulation and visualization
import matplotlib.pyplot as plt
from PIL import Image
import seaborn as sns
import pandas as pd
import numpy as np

# Deep Learning libraries
import torch
import torchvision
import torchsummary
from torch.utils import data
from torchvision import datasets, models, transforms

# Set seed for reproducibility
SEED = 42
np.random.seed(SEED)

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

from torch.utils.data import DataLoader, Dataset
from glob import glob
from itertools import chain
import random

class EuroSATDataset(Dataset):

    def __init__(self, root, train_flag = True):

        self.train_flag = train_flag
        self.root = root
        self.images_paths = [glob(f'{root}/{folder}/*.jpg') for folder
in os.listdir(f"{root}")]
        self.images_paths =
list(chain.from_iterable(self.images_paths))
        random.shuffle(self.images_paths)

        # комментарий преподавателя: на текущей машине очень мало
мощностей, максимально обрезаем данные
        # для обучения и валидации в соотношении 100/30

        self.count = {True: 100}
        self.count[False] = 30

```

```

        self.classes_names = {class_name:label for label, class_name
in enumerate(os.listdir(f"{root}"))}
        self.labels =
[self.classes_names[os.path.basename(os.path.dirname(path))] for path
in self.images_paths]

        # применяем аугментациб данных. В качестве ДЗ расширите список
применяемых преобразований.
        self.transform_train = transforms.Compose([
            transforms.RandomResizedCrop((224, 224)),
            transforms.RandomHorizontalFlip(),
            transforms.RandomRotation(30),
            transforms.ColorJitter(brightness=0.2, contrast=0.2),
            transforms.ToTensor(),
            transforms.Normalize(mean=[0.485, 0.456, 0.406],
std=[0.229, 0.224, 0.225])
        ])

        self.transform_test = transforms.Compose([
            transforms.Resize((224, 224)),
            transforms.ToTensor(),
            transforms.Normalize(mean=[0.485, 0.456, 0.406],
std=[0.229, 0.224, 0.225])
        ])

    def __len__(self):
        return self.count[self.train_flag]

    def __getitem__(self, index):
        if (self.train_flag):
            index = index % self.count[True]
        else:
            index = self.count[True] + index % self.count[False]
        image_path = self.images_paths[index]
        image = Image.open(image_path).convert('RGB')
        label = self.labels[index]

        if self.train_flag:
            image = self.transform_train(image)
        else:
            image = self.transform_test(image)

        return image.float().to(device),
torch.tensor([label]).float().to(device)

data = EuroSATDataset('./EuroSAT/2750', train_flag=True)
len(data)

import random

```

```

import cv2

image, label = data[random.randint(0, len(data))]
print(f"Image Size: {image.shape[2]} x {image.shape[1]} x {image.shape[0]}")
print(f"Label: {label}")
print([key for key, value in data.classes_names.items() if value == label][0])
plt.imshow(image.permute(1,2,0).cpu().numpy())
plt.show()

train_dataset = EuroSATDataset('./EuroSAT/2750')
test_dataset = EuroSATDataset('./EuroSAT/2750', train_flag = False)
train_dataset_loader = DataLoader(train_dataset, batch_size=32,
shuffle=True, drop_last=True)
test_dataset_loader = DataLoader(test_dataset, batch_size=1,
shuffle=True, drop_last=True)

# fine tuning MODEL
# change last model layer (model.fc) for classification on EuroSAT set
of classes, 10 classes
# (not 1000 like for ImageNet dataset pretrained weights)
# pretrained models https://pytorch.org/vision/stable/models.html

model = models.resnet50(weights=models.ResNet50_Weights.DEFAULT)
model.fc = torch.nn.Linear(model.fc.in_features, 10)
model = model.to(device)
torchsummary.summary(model, (3, 224, 224))

# Specify number of epochs and learning rate
lr = 1e-3

# Specify criterion and optimizer
criterion = torch.nn.CrossEntropyLoss()
optimizer = torch.optim.SGD(model.parameters(), lr=lr)

def train_step(model, loss_function, optimizer, image, label):
    model.train()
    optimizer.zero_grad()
    prediction = model(image)
    #print(prediction)
    #print(label)
    loss = loss_function(prediction.squeeze(), label.long().squeeze())
    loss.backward()
    optimizer.step()
    return loss.item()

@torch.no_grad()
def accuracy(model, loss_function, image, label):
    model.eval()

```

```

prediction = model(image)
max_values, argmaxes = prediction.max(-1)
is_correct = argmaxes == label.int().squeeze()
return is_correct.cpu().numpy().tolist()

@torch.no_grad()
def validation_loss(model, loss_function, image, label):
    model.eval()
    prediction = model(image)
    loss = loss_function(prediction.squeeze(), label.long().squeeze())
    return loss.item()

train_losses = []
train_accuracies = []
test_losses = []
test_accuracies = []

n_epochs = 10
for epoch in range(n_epochs):
    print(f"Epoch: {epoch+1}")
    train_epoch_losses = []
    train_epoch_accuracies = []
    for idx, (image, label) in enumerate(train_dataset_loader):
        loss = train_step(model, criterion, optimizer, image, label)
        train_epoch_losses.append(loss)
        print(loss)
        #if (idx + 1) %100 == 0: print(loss)
    train_epoch_loss = np.mean(train_epoch_losses)
    print(f"Train Loss: {train_epoch_loss:.4f}")
    train_losses.append(train_epoch_loss)

    for idx, (image, label) in enumerate(train_dataset_loader):
        is_correct = accuracy(model, criterion, image, label)
        train_epoch_accuracies.extend(is_correct)
    train_epoch_accuracy = np.mean(train_epoch_accuracies)
    print(f"Train Accuracy: {train_epoch_accuracy*100:.2f}%")
    train_accuracies.append(train_epoch_accuracy)

    test_epoch_losses = []
    test_epoch_accuracies = []
    for idx, (image, label) in enumerate(test_dataset_loader):
        #print(label)
        loss = validation_loss(model, criterion, image, label)
        test_epoch_losses.append(loss)
        is_correct = accuracy(model, criterion, image, label)
        test_epoch_accuracies.extend(is_correct)
    test_epoch_loss = np.mean(test_epoch_losses)
    print(f"Test Loss: {test_epoch_loss:.4f}")
    test_losses.append(test_epoch_loss)
    test_epoch_accuracy = np.mean(test_epoch_accuracies)

```

```

print(f"Test Accuracy: {test_epcoh_accuracy*100:.2f}%")
test accuracies.append(test_epcoh_accuracy)

# example of model saving
model_dir = "./models/"
if not os.path.exists(model_dir):
    os.makedirs(model_dir)

model_file = os.path.join(model_dir, 'best_model.pth')
model_file
torch.save(model.state_dict(), model_file)


# EXAMPLE OF MODEL LOADING
model = models.resnet50(weights=models.ResNet50_Weights.DEFAULT)
model.fc = torch.nn.Linear(model.fc.in_features, 10)
model = model.to(device)

# example of loading model from the file
model.load_state_dict(torch.load(model_file, weights_only=True))

Clipping input data to the valid range for imshow with RGB data
([0..1] for floats or [0..255] for integers). Got range [-
2.117904..0.23477131].

Image Size: 224 x 224 x 3
Label: tensor([2.])
HerbaceousVegetation

```



Downloading: "https://download.pytorch.org/models/resnet50-11ad3fa6.pth" to C:\Users\Даниил\.cache\torch\hub\checkpoints\resnet50-11ad3fa6.pth  
100%|

| 97.8M/97.8M [00:42<00:00, 2.43MB/s]

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 64, 112, 112]	9,408
BatchNorm2d-2	[-1, 64, 112, 112]	128
ReLU-3	[-1, 64, 112, 112]	0
MaxPool2d-4	[-1, 64, 56, 56]	0
Conv2d-5	[-1, 64, 56, 56]	4,096
BatchNorm2d-6	[-1, 64, 56, 56]	128
ReLU-7	[-1, 64, 56, 56]	0
Conv2d-8	[-1, 64, 56, 56]	36,864
BatchNorm2d-9	[-1, 64, 56, 56]	128
ReLU-10	[-1, 64, 56, 56]	0
Conv2d-11	[-1, 256, 56, 56]	16,384
BatchNorm2d-12	[-1, 256, 56, 56]	512
Conv2d-13	[-1, 256, 56, 56]	16,384
BatchNorm2d-14	[-1, 256, 56, 56]	512
ReLU-15	[-1, 256, 56, 56]	0

Bottleneck-16	[-1, 256, 56, 56]	0
Conv2d-17	[-1, 64, 56, 56]	16,384
BatchNorm2d-18	[-1, 64, 56, 56]	128
ReLU-19	[-1, 64, 56, 56]	0
Conv2d-20	[-1, 64, 56, 56]	36,864
BatchNorm2d-21	[-1, 64, 56, 56]	128
ReLU-22	[-1, 64, 56, 56]	0
Conv2d-23	[-1, 256, 56, 56]	16,384
BatchNorm2d-24	[-1, 256, 56, 56]	512
ReLU-25	[-1, 256, 56, 56]	0
Bottleneck-26	[-1, 256, 56, 56]	0
Conv2d-27	[-1, 64, 56, 56]	16,384
BatchNorm2d-28	[-1, 64, 56, 56]	128
ReLU-29	[-1, 64, 56, 56]	0
Conv2d-30	[-1, 64, 56, 56]	36,864
BatchNorm2d-31	[-1, 64, 56, 56]	128
ReLU-32	[-1, 64, 56, 56]	0
Conv2d-33	[-1, 256, 56, 56]	16,384
BatchNorm2d-34	[-1, 256, 56, 56]	512
ReLU-35	[-1, 256, 56, 56]	0
Bottleneck-36	[-1, 256, 56, 56]	0
Conv2d-37	[-1, 128, 56, 56]	32,768
BatchNorm2d-38	[-1, 128, 56, 56]	256
ReLU-39	[-1, 128, 56, 56]	0
Conv2d-40	[-1, 128, 28, 28]	147,456
BatchNorm2d-41	[-1, 128, 28, 28]	256
ReLU-42	[-1, 128, 28, 28]	0
Conv2d-43	[-1, 512, 28, 28]	65,536
BatchNorm2d-44	[-1, 512, 28, 28]	1,024
Conv2d-45	[-1, 512, 28, 28]	131,072
BatchNorm2d-46	[-1, 512, 28, 28]	1,024
ReLU-47	[-1, 512, 28, 28]	0
Bottleneck-48	[-1, 512, 28, 28]	0
Conv2d-49	[-1, 128, 28, 28]	65,536
BatchNorm2d-50	[-1, 128, 28, 28]	256
ReLU-51	[-1, 128, 28, 28]	0
Conv2d-52	[-1, 128, 28, 28]	147,456
BatchNorm2d-53	[-1, 128, 28, 28]	256
ReLU-54	[-1, 128, 28, 28]	0
Conv2d-55	[-1, 512, 28, 28]	65,536
BatchNorm2d-56	[-1, 512, 28, 28]	1,024
ReLU-57	[-1, 512, 28, 28]	0
Bottleneck-58	[-1, 512, 28, 28]	0
Conv2d-59	[-1, 128, 28, 28]	65,536
BatchNorm2d-60	[-1, 128, 28, 28]	256
ReLU-61	[-1, 128, 28, 28]	0
Conv2d-62	[-1, 128, 28, 28]	147,456
BatchNorm2d-63	[-1, 128, 28, 28]	256
ReLU-64	[-1, 128, 28, 28]	0

Conv2d-65	[-1, 512, 28, 28]	65,536
BatchNorm2d-66	[-1, 512, 28, 28]	1,024
ReLU-67	[-1, 512, 28, 28]	0
Bottleneck-68	[-1, 512, 28, 28]	0
Conv2d-69	[-1, 128, 28, 28]	65,536
BatchNorm2d-70	[-1, 128, 28, 28]	256
ReLU-71	[-1, 128, 28, 28]	0
Conv2d-72	[-1, 128, 28, 28]	147,456
BatchNorm2d-73	[-1, 128, 28, 28]	256
ReLU-74	[-1, 128, 28, 28]	0
Conv2d-75	[-1, 512, 28, 28]	65,536
BatchNorm2d-76	[-1, 512, 28, 28]	1,024
ReLU-77	[-1, 512, 28, 28]	0
Bottleneck-78	[-1, 512, 28, 28]	0
Conv2d-79	[-1, 256, 28, 28]	131,072
BatchNorm2d-80	[-1, 256, 28, 28]	512
ReLU-81	[-1, 256, 28, 28]	0
Conv2d-82	[-1, 256, 14, 14]	589,824
BatchNorm2d-83	[-1, 256, 14, 14]	512
ReLU-84	[-1, 256, 14, 14]	0
Conv2d-85	[-1, 1024, 14, 14]	262,144
BatchNorm2d-86	[-1, 1024, 14, 14]	2,048
Conv2d-87	[-1, 1024, 14, 14]	524,288
BatchNorm2d-88	[-1, 1024, 14, 14]	2,048
ReLU-89	[-1, 1024, 14, 14]	0
Bottleneck-90	[-1, 1024, 14, 14]	0
Conv2d-91	[-1, 256, 14, 14]	262,144
BatchNorm2d-92	[-1, 256, 14, 14]	512
ReLU-93	[-1, 256, 14, 14]	0
Conv2d-94	[-1, 256, 14, 14]	589,824
BatchNorm2d-95	[-1, 256, 14, 14]	512
ReLU-96	[-1, 256, 14, 14]	0
Conv2d-97	[-1, 1024, 14, 14]	262,144
BatchNorm2d-98	[-1, 1024, 14, 14]	2,048
ReLU-99	[-1, 1024, 14, 14]	0
Bottleneck-100	[-1, 1024, 14, 14]	0
Conv2d-101	[-1, 256, 14, 14]	262,144
BatchNorm2d-102	[-1, 256, 14, 14]	512
ReLU-103	[-1, 256, 14, 14]	0
Conv2d-104	[-1, 256, 14, 14]	589,824
BatchNorm2d-105	[-1, 256, 14, 14]	512
ReLU-106	[-1, 256, 14, 14]	0
Conv2d-107	[-1, 1024, 14, 14]	262,144
BatchNorm2d-108	[-1, 1024, 14, 14]	2,048
ReLU-109	[-1, 1024, 14, 14]	0
Bottleneck-110	[-1, 1024, 14, 14]	0
Conv2d-111	[-1, 256, 14, 14]	262,144
BatchNorm2d-112	[-1, 256, 14, 14]	512
ReLU-113	[-1, 256, 14, 14]	0



Conv2d-114	[-1, 256, 14, 14]	589,824
BatchNorm2d-115	[-1, 256, 14, 14]	512
ReLU-116	[-1, 256, 14, 14]	0
Conv2d-117	[-1, 1024, 14, 14]	262,144
BatchNorm2d-118	[-1, 1024, 14, 14]	2,048
ReLU-119	[-1, 1024, 14, 14]	0
Bottleneck-120	[-1, 1024, 14, 14]	0
Conv2d-121	[-1, 256, 14, 14]	262,144
BatchNorm2d-122	[-1, 256, 14, 14]	512
ReLU-123	[-1, 256, 14, 14]	0
Conv2d-124	[-1, 256, 14, 14]	589,824
BatchNorm2d-125	[-1, 256, 14, 14]	512
ReLU-126	[-1, 256, 14, 14]	0
Conv2d-127	[-1, 1024, 14, 14]	262,144
BatchNorm2d-128	[-1, 1024, 14, 14]	2,048
ReLU-129	[-1, 1024, 14, 14]	0
Bottleneck-130	[-1, 1024, 14, 14]	0
Conv2d-131	[-1, 256, 14, 14]	262,144
BatchNorm2d-132	[-1, 256, 14, 14]	512
ReLU-133	[-1, 256, 14, 14]	0
Conv2d-134	[-1, 256, 14, 14]	589,824
BatchNorm2d-135	[-1, 256, 14, 14]	512
ReLU-136	[-1, 256, 14, 14]	0
Conv2d-137	[-1, 1024, 14, 14]	262,144
BatchNorm2d-138	[-1, 1024, 14, 14]	2,048
ReLU-139	[-1, 1024, 14, 14]	0
Bottleneck-140	[-1, 1024, 14, 14]	0
Conv2d-141	[-1, 512, 14, 14]	524,288
BatchNorm2d-142	[-1, 512, 14, 14]	1,024
ReLU-143	[-1, 512, 14, 14]	0
Conv2d-144	[-1, 512, 7, 7]	2,359,296
BatchNorm2d-145	[-1, 512, 7, 7]	1,024
ReLU-146	[-1, 512, 7, 7]	0
Conv2d-147	[-1, 2048, 7, 7]	1,048,576
BatchNorm2d-148	[-1, 2048, 7, 7]	4,096
Conv2d-149	[-1, 2048, 7, 7]	2,097,152
BatchNorm2d-150	[-1, 2048, 7, 7]	4,096
ReLU-151	[-1, 2048, 7, 7]	0
Bottleneck-152	[-1, 2048, 7, 7]	0
Conv2d-153	[-1, 512, 7, 7]	1,048,576
BatchNorm2d-154	[-1, 512, 7, 7]	1,024
ReLU-155	[-1, 512, 7, 7]	0
Conv2d-156	[-1, 512, 7, 7]	2,359,296
BatchNorm2d-157	[-1, 512, 7, 7]	1,024
ReLU-158	[-1, 512, 7, 7]	0
Conv2d-159	[-1, 2048, 7, 7]	1,048,576
BatchNorm2d-160	[-1, 2048, 7, 7]	4,096
ReLU-161	[-1, 2048, 7, 7]	0
Bottleneck-162	[-1, 2048, 7, 7]	0

Conv2d-163	[-1, 512, 7, 7]	1,048,576
BatchNorm2d-164	[-1, 512, 7, 7]	1,024
ReLU-165	[-1, 512, 7, 7]	0
Conv2d-166	[-1, 512, 7, 7]	2,359,296
BatchNorm2d-167	[-1, 512, 7, 7]	1,024
ReLU-168	[-1, 512, 7, 7]	0
Conv2d-169	[-1, 2048, 7, 7]	1,048,576
BatchNorm2d-170	[-1, 2048, 7, 7]	4,096
ReLU-171	[-1, 2048, 7, 7]	0
Bottleneck-172	[-1, 2048, 7, 7]	0
AdaptiveAvgPool2d-173	[-1, 2048, 1, 1]	0
Linear-174	[-1, 10]	20,490

```

=====
Total params: 23,528,522
Trainable params: 23,528,522
Non-trainable params: 0
-----

```

```

Input size (MB): 0.57
Forward/backward pass size (MB): 286.55
Params size (MB): 89.75
Estimated Total Size (MB): 376.88
-----

```

```

Epoch: 1
2.2946979999542236
2.3574068546295166
2.3407657146453857
Train Loss: 2.3310
Train Accuracy: 4.17%
Test Loss: 2.3344
Test Accuracy: 6.67%
Epoch: 2
2.377584934234619
2.3298561573028564
2.286466360092163
Train Loss: 2.3313
Train Accuracy: 8.33%
Test Loss: 2.3362
Test Accuracy: 6.67%
Epoch: 3
2.317997694015503
2.3509457111358643
2.311251640319824
Train Loss: 2.3267
Train Accuracy: 10.42%
Test Loss: 2.3207
Test Accuracy: 10.00%
Epoch: 4
2.339057683944702
2.3189327716827393

```

2.2845723628997803  
Train Loss: 2.3142  
Train Accuracy: 9.38%  
Test Loss: 2.3078  
Test Accuracy: 6.67%  
Epoch: 5  
2.302266836166382  
2.3094396591186523  
2.3284969329833984  
Train Loss: 2.3134  
Train Accuracy: 8.33%  
Test Loss: 2.3024  
Test Accuracy: 0.00%  
Epoch: 6  
2.2890167236328125  
2.3025779724121094  
2.3293771743774414  
Train Loss: 2.3070  
Train Accuracy: 6.25%  
Test Loss: 2.3052  
Test Accuracy: 0.00%  
Epoch: 7  
2.3336730003356934  
2.313509941101074  
2.3204541206359863  
Train Loss: 2.3225  
Train Accuracy: 12.50%  
Test Loss: 2.3135  
Test Accuracy: 3.33%  
Epoch: 8  
2.320241689682007  
2.298161745071411  
2.3121495246887207  
Train Loss: 2.3102  
Train Accuracy: 7.29%  
Test Loss: 2.3183  
Test Accuracy: 3.33%  
Epoch: 9  
2.305762529373169  
2.2849740982055664  
2.281937599182129  
Train Loss: 2.2909  
Train Accuracy: 7.29%  
Test Loss: 2.3182  
Test Accuracy: 6.67%  
Epoch: 10  
2.253018379211426  
2.3571977615356445  
2.275646448135376  
Train Loss: 2.2953

Train Accuracy: 10.42%

Test Loss: 2.3139

Test Accuracy: 6.67%

<All keys matched successfully>